

COMPASS → PEER — tenure dataset bolstering (scrape + adjudicate)

Last synced: 2026-09-01

Audience: Charles + PEER (working session). COMPASS planning handoff — not for Alex yet.

Standalone: Alex wants HERO / F-HERO-style metrics on tenure soon; the panel is not robust enough yet. This doc is the ordered checklist to expand scrape coverage before any tenure hero work.

Deep reference: `tenure/documents/TENURE_SCAPE_AND_ADJUDICATE_FOR_DUMMIES.md` · `scripts/DATA_SYNC.md`

YOU ARE HERE

Step	Task	Owner	Status
0	Mac/Rivanna/git synced (<code>0d1802a</code>)	Charles	✓
1	Pull latest tenure outputs to Mac	Charles	→ NEXT
2	Review audit + pick expansion lane(s) A–E	Charles + PEER	Pending
3	Adjudicate URLs / add schools (Phase 1)	Charles + PEER	Pending
4	Commit worksheet changes + push/rsync to Rivanna	Charles	Pending
5	Run scrape chain (Cells 2–4) on Rivanna	PEER or Charles	Pending
6	Morning check + rsync pull to Mac	Charles	Pending
7	Rebuild panel (Cells 5–9) if scrape healthy	PEER or Charles	Pending
8	Lock tenure HERO spec with Alex	Charles + Alex	After 7

Alex one-liner (context only): Explore hero / F-HERO on tenure — **after** the faculty panel is worth plotting.

Why now (30 seconds)

MBB Wang-arc work is in good shape on the sports side (empirical HERO/F-HERO, reigning-hero calibration, sim replay on 2009–21). The **next domain** is tenure. The pipeline code exists; the **corpus** (HTML coverage, URL quality, panel rows) needs a deliberate expansion round before hero metrics mean anything.

Reading docs does **not** add HTML. New data = **Phase 1 (Mac inputs) + Phase 2 (Rivanna Wayback run)**.

The two-phase loop (memorize this)

Phase 1 – Mac: edit school list / URL worksheet → apply → git push + rsync push
Phase 2 – Rivanna: `540_tenure_pipeline.ipynb` Cell 0 → `sbatch pipe_job.slurm`

After Phase 2: **rsync pull tenure** to Mac so Cursor sees new JSONL/logs.

Expansion lanes — pick with PEER

Lane	You want...	Mac edit	Rivanna flags (minimum)
A	New R1 CS schools	<code>r1_cs_departments.csv</code> , <code>r1_schools_data.py</code>	CELL2, 3A CDX, 3B download, 4
B	Better URL for weak school	<code>url_update_worksheet.csv</code> → <code>apply_url_updates.py</code>	Same as A
C	URLs never CDX-queried	Often none	3A CDX (+ download, 4)
D	CDX done, HTML missing	None	3A CDX=False , download, 4
E	Prior CDX timeouts	None	3 retry + download, 4

Most common today: **B** (fix weak parsers) and/or **A** (grow toward full R1 list).

Optional URL leads:

```
python tenure/tenure_pipeline/discover_faculty_urls.py
# → faculty_url_suggestions.csv → paste into new_url → apply_url_updates.py
```

Itemized checklist

Before the PEER session (Charles — ~30 min)

- ☐ 1. From repo root on Mac:

```
./scripts/rsync_pull_recent_hpc.sh tenure
```

- ☐ 2. Open and skim (share screen with PEER):

- `tenure/tenure_pipeline/faculty_snapshots_strategy_audit.jsonl`
- `tenure/tenure_pipeline/url_update_worksheet.csv`
- `faculty_url_suggestions.csv` (if present after pull)

- ☐ 3. Have `TENURE_SCRAPING_AND_ADJUDICATE_FOR_DUMMIES.md` open — PEER does not need the 900-line overview.

During the PEER session — decide together

- ☐ 4. **Scope:** How many schools this round? (fix worst parsers only vs add N new departments)
- ☐ 5. **Lane(s):** Mark A / B / C / D / E from table above
- ☐ 6. **Split labor:**
 - Who edits `new_url` rows?
 - Who adds new schools to CSV?
 - Who submits Slurm on Rivanna?
- ☐ 7. **Quality bar:** What counts as “good enough” to run panel rebuild? (e.g. audit red schools cleared, min HTML count per school)
- ☐ 8. **OpenAlex this round?** Cell 6A/6B now, or scrape/HTML first and OA later?

Mac — Phase 1 (after decisions)

- ☐ 9. Edit worksheets / school list per lane A or B
- ☐ 10. If URLs changed:

```
python tenure/tenure_pipeline/apply_url_updates.py
```

- ☐ 11. Stage + commit worksheet/school-list changes (small, focused commit)
- ☐ 12. Push to GitHub + push code to Rivanna:

```
git push origin main
./scripts/rsync_push_to_hpc.sh
```

Rivanna — Phase 2 (PEER or Charles)

- ☐ 13. Confirm clone current:

```
cd ~/Ivy_Net && git pull origin main
```

- ☐ 14. Open `tenure/540_tenure_pipeline.ipynb` — Cell 0 for **new scrape minimum**:
 - `RUN_CELL2`, `RUN_CELL3_CDX`, `RUN_CELL3_DOWNLOAD`, `RUN_CELL4` = **True**
 - All other `RUN_CELL*` = **False** (unless step 8 said otherwise)
- ☐ 15. Submit:

```
cd ~/Ivy_Net && sbatch pipe_job.slurm
queue -u dzk3ja
```

- ☐ 16. Next morning — read log:

```
tail -80 ~/Ivy_Net/slurm_out/slurm-pipe_job-*.out
```

Look for: ERROR lines, download counts, parse summary

After a healthy scrape — Phase 3

- ☐ 17. Mac pull:

```
./scripts/rsync_pull_recent_hpc.sh tenure
```

- ☐ **18. Workflow E** (panel + LOO metrics + plots) — Cell 0:
 - Cell 5 → `faculty_panel.jsonl`
 - Cells 7–9 → enriched panel, pools, inverted-U plot
- ☐ **19.** Sanity before Alex / tenure HERO:
 - Panel row count ↑ vs last run?
 - Weak schools improved in audit?
 - Ready to lock **Y**, **Â**, LOO pool definition (separate Alex memo — not this doc)

Who owns what (default split)

Piece	Mac	Rivanna
URL research + worksheet	Charles + PEER	—
<code>apply_url_updates.py</code> , git commit	Charles	—
Wayback CDX / download / parse	—	PEER or Charles
Slurm submit / log triage	—	PEER or Charles
rsync pull for Cursor	Charles	—
Panel rebuild (Cells 5–9)	trigger from either	runs on HPC

Git = code + small CSVs. **rsync** = JSONL, HTML tree, slurm logs. See `DATA_SYNC.md`.

Success criteria (this round)

- Chosen expansion lane(s) executed without Slurm hard failure
- Audit shows material improvement for targeted weak schools **or** new schools have first HTML + parse rows
- Mac has fresh `tenure/tenure_pipeline/` outputs via rsync
- Charles can answer: “How many person-years in panel? How many schools at min coverage?”

After bolstering — what COMPASS needs from Charles + Alex

Not for PEER session. After step 19, open a short **tenure HERO lock** with Alex:

- Outcome **Y** (tenure granted? promotion event?)
- Panel aperture (years, schools, activity filters)
- Â** field(s) for talent / F-HERO slices
- LOO pool definition for tenure HERO
- Empirical-only first vs generative replay (MBB template)

COMPASS turns that into scripts and sandbox paths.

Thread log

Date	Entry
2026-09-01	COMPASS draft after Mac git recovery; Alex tenure HERO interest; PEER scrape round queued. MBB at <code>0d1802a</code> ; Rivanna pulled same.