

The Big Fish and the Small Pond — A Porch Reading for Charles

Prepared by COMPASS (with SCOUT’s numbers and cross-read folded in)
Date: August 20, 2026
SCOUT review: Complete (2026-08-21) — see `20260821_SCOUT_porch_reading_review.md`. Three minor clarifications only; approved for porch use.
Purpose: One coherent story you can print, bring home, and read quietly on the porch with Ginger — mission, contention, what we learned today, and where we go next.
For Ginger: You don’t need the notation. The Squid-and-Jackal story is: *same talent, different team — does standing out on the roster change who gets drafted?* Charles has been chasing that question for years. Tonight’s read is where the data stand.
Source memos merged: COMPASS BDP wisdom · SCOUT+COMPASS joint CCT wisdom · SCOUT §11 addendum · Track C robustness note.
This document replaces reading three separate agent memos. The technical pair stays in the folder for reference.

A note before you begin

You asked SCOUT and COMPASS to look at your new **Basic Data Plots** deck, connect it to **HAND17 Slide 5** (team talent intervals), and help you find how to **illuminate** a conviction you have carried for a long time: the **big fish / small pond** story.
We read everything — all 21 slide exports, every BDP JSON sidecar, the interval overlap slide, the hero history, the model memos, and the BPM/OBPM robustness work.
Good news first: Your Central Contention and Theme (CCT) is **coherent**. It fits the Army evidence, the model structure, the interval geometry, and the new BDP overlays. Nothing we saw today falsifies it.
Honest news second: The CCT is **not yet visible** in the charts you have been staring at the longest — the marginal “hero” curve and the new histograms. Those charts answer **different questions**. The Squid-versus-Jackal test needs a **microscope**, not a wider lens.
SCOUT and COMPASS agree on that. There are **no disagreements** between us on substance today. Where one of us said something exclusively, you will see **[SCOUT]** or **[COMPASS]** in the margin.

Part 1 — What you told us (in your words)

You said something like this:

*I feel like if we go to a pretty good team — call them the **Squids** — there will be a great player. I feel like if we go to a contender — call them the **Jackals** — there will be **several** players with the **same talent** as that one great Squid player. And I think the Squid star, considering everything, will empirically be **more likely to be drafted** than each of the corresponding players at his level on the Jackals — because of **congestion**.*

You are convinced this is true. In Army data it is **extreme**. In NCAA data you suspect it is a **subtle nuance** — something you have to **zoom in** on and find.
You also said:

- You built the **BDP deck** to get into the **nitty gritty** of what happens among ****top T_j teams**** — intervals, congestion, roster sizes — not just the league-wide marginals.
- You remember **HAND17 Slide 5**: teams differ in T_j, intervals **move up** as teams get better, and overlap/stacking at a given talent level matters for your story.
- You noticed something in the new slides: **OBPM** and **BPM** DFT populations look **tighter** (smaller standard deviations) than the non-DFT populations — and you wondered if that has implications.
- You want populations, filters, and scales that **bring the CCT to empirical day** — not another cryptic outline, but a **view of the data** that makes the principle visible even if NCAA whispers.

We heard you. This document is our answer.

Part 2 — The Central Contention, in plain English

The question (this is NOT “do good players get drafted?”)
Everyone knows good players get drafted more often. That is not the dissertation move.
The move is:
Hold individual talent fixed. Then ask: does **where you sit on the team** — how **crowded** your pond is with other strong players — change your odds of being drafted?

The cast

Name	Who they are	Role in the story
Squid star	Best player on a pretty good team	“Big fish, small pond” — stands out locally
Jackal peer	Player of similar individual talent on a contender	“Small fish, big pond” — same fish, more rivals at that level
Congestion	Too many viable peers on the roster	Hurts you in the race for scarce draft attention , not necessarily in your box score

What “congestion” is NOT (binding rule — [COMPASS])
Congestion in your **draft story** is **not** the same thing as “bad teammates hurt my development.”

- The dissertation deliberately separates:
1. **Environment** — peers can help you grow (B) or crowd you in development (D). Net environment is $L_{net} \approx B - D$.
 2. **Score** — how we **rank** candidates: $S_i = talent - \lambda \times L_C$ (congestion in the **rank**).
 3. **Select** — who **wins** the scarce slots **after** ranking (top-K, Gibbs, etc.).
 4. **Hero** — what we **observe** in real data (draft rate vs teammate quality bins).

Do not merge these. The hero is an **outcome**. It is not the same object as S_i , and environment is not advancement.

Army vs NCAA

	Army	NCAA (your working guess)
Signal strength	Screams — obvious, robust tail drop even at macro level	Whispers — middle rise, elite tail flat in marginals
Baseline outcome rate	~35–40% of Captains promote — high baseline, mostly an assembly-line process	~2–2.5% drafted on your panel — scarcer binary outcome
Why it screams (not K/N)	Peer/environment effects show up strongly in the aggregate hero ; survival analysis adds timing and subtle peer indicators within a process where many eventually succeed	Same CCT logic may be present, but marginals + overlap + rare draft bury the Squid–Jackal contrast until you condition
What you need	You already had the macro shape	A microscope — matched $\hat{A} \times$ pond plots

Do not confuse “Army screams” with “Army K/N is tiny.” Captain promotion is **not** a 1%-style lottery. The Army is loud because the **tail drop is real and robust in the display** — not because slots are scarcer than the NBA draft (they are not).

[SCOUT]: Do not expect another marginal ventile curve to “dip at the elite bin” on POST-QC NCAA data. If CCT is true in NCAA, it will show up when you **hold $\hat{A}$ fixed** and compare ponds — not in a single averaged line.

Part 3 — Where this sits in the bigger dissertation story

You started in the **Army**: promotion and selection in peer environments. You found a pattern — advancement vs peer quality often **risers through the middle** and sometimes **softens at the very top** — a **robust macro-level tail drop** that does not need fancy conditioning to see. Most Captains are on a similar track (~35–40% promote); **survival analysis** then teases **when** promotion happens and which subtle peer indicators matter inside that mostly-uniform assembly line. You asked: is that special to the Army?

Basketball became the working example because you have rich data, a clear scarce outcome (NBA draft), and one slide that captures the phenomenon — the **hero**.

The **Wang-style ladder** you are climbing:

1. **Phenomenon (Layer A)** — What does the data look like? (*Hero, BDP marginals*)
2. **Mechanism story (Layer B)** — Why might peers help and hurt? Congestion in the **score**, not merged with environment. (*Words first*)
3. **Minimal generative test (Layer C)** — Does the smallest fake league show the mechanism **can** bend selection? (*Phase B sim, λ ablation*)
4. **Predictions / fit (Rung 3+)** — MLE, PD14-style checks, conditional empirics. (*Where CCT lives next*)

[COMPASS]: Act I of the congestion chapter is now the **BDP deck** you finished today. Act II is **conditional** draft-rate geometry — not another histogram.

Part 4 — What you just built (BDP deck) and why it matters

You finished `Basics_data_plots_HAND.pptx` — 21 slides. Source PNGs live in `basic_data_plots/`. The deck walks four **base populations** (all 2011–2021, QC = dash placeholder names removed):

Population	Plain English
FP 11_21	Full panel — no minimum games, no minutes floor. Messy but honest about raw ESPN.
mg10 min0	Drop fragmentary team-seasons (≤ 10 games). Still all minute levels.
mg10 min10	Same, but players need ≥ 10 minutes that season.
mg10 min20	Same, but ≥ 20 minutes — this is your locked hero / ASSIGN panel .

For each population (except roster-only slides), you show:

- $\hat{A}_i$ and T_i distributions (PPM, then BPM, then OBPM)

- Blue bars = **without DFT** (full filtered population)
- Orange line = **+DFT** (teams that **ever** had a draftee in the 2011–2021 window)
- Roster size $|T_j|$ at the end

[COMPASS]: The deck is **Act I: population geometry**. It answers: *What does the league look like? What do draft-ecosystem teams look like compared to the league?* It does **not** yet answer: *At the same $\hat{A}$, does the Squid beat the Jackal?*
That distinction is the whole pivot of this porch reading.

Part 5 — What the new slides actually show (numbers in plain English)

5.1 Players spread wide; team averages spread narrow

On almost every BDP slide:

- **Individual player talent $\hat{A}_i$** has standard deviation ≈ 1.0 (by design — within-season z-scores).
- ****Team mean talent T_j **** has standard deviation ≈ 0.25 to **0.43**, depending on metric and filter.

Plain English: Stars and bench players differ a lot. **Average team strength differs much less**. Teams are not identical — but team *means* live in a **tighter band** than players do.

This matches HAND17 Slide 5: every team has a **wide internal window** of talent (~ 3 z from weakest to strongest player), but when you sort teams by T_j , the **whole window marches upward** for contenders.

5.2 The orange +DFT line — draft-ecosystem zoom

When you overlay teams and players from programs that **ever produced a draft pick** in the window:

At mg10 min20 (your hero panel), PPM example:

- **Players (+DFT):** mean shifts up ($\sim +0.09$ z); spread **narrows** (sd 1.00 $\rightarrow$ 0.76, about **24% tighter**).
- **Teams (+DFT):** mean shifts up ($\sim +0.09$ z); spread **narrows** again (T_j sd 0.25 $\rightarrow$ 0.19, about **26% tighter**).
- **Counts:** $\sim 46,000$ player-seasons without DFT vs $\sim 17,700$ +DFT; $\sim 3,842$ team-seasons vs $\sim 1,494$ +DFT.

Plain English: Draft-history teams are **better on average** and **more alike each other** than the full league. The Jackals (contender-ish teams) exist — they are the **right tail** of T_j — but they are **few** and **clustered**. You will not see congestion by staring at a histogram of all teams. You must **condition**: pick a talent level, then compare ponds.

5.3 What you noticed — OBPM vs BPM compression [SCOUT]

At **mg10 min20**, when you compare full league to +DFT, **team-level** spread shrinks more under BPM-family metrics than PPM:

Metric	T_j sd (all)	T_j sd (+DFT)	How much narrower?
PPM	0.251	0.186	26%
BPM	0.580	0.314	46%
OBPM	0.428	0.245	43%

[SCOUT]’s read: BPM and OBPM absorb **team role and context**, not just raw production. When you restrict to draft-ecosystem teams, comprehensive metrics **compress** — as if everyone on those rosters gets **re-scaled into a tighter impact band**. At the **team** level, **OBPM** DFT teams are **more alike** (sd 0.25) than **BPM** DFT teams (sd 0.31) — plausible if defensive/role noise widens total BPM spreads.

[COMPASS]: Treat BPM/OBPM as **magnifying glasses**, not a new canonical hero. Track C already showed: OBPM/BPM **do not restore** the July elite-bin dip on POST-QC data; they show a **stronger monotonic rise** in the marginal hero curve. Footnote always: BPM merge drops $\sim 1,150$ player-seasons vs PPM (Sports-Reference coverage).

5.4 Roster count is probably not the story [SCOUT]

Team-size slides:

- At **FP** (no minutes floor): draft-history teams average **slightly fewer** players (15.4 vs 16.8) — junk roster inflation lives in non-DFT teams.
- At **mg10 min10** and **min20**: DFT and non-DFT are **virtually the same** (~ 12 players with real minutes; mean $|T_j| \approx 12.0$ vs 11.8 at min20; max 19 both).

Plain English: Kentucky and Mid-Major U. do **not** differ much on “how many guys played meaningful minutes.” Congestion in your story is **talent stacked into similar roster slots**, not fewer names on the sheet.

5.5 Which population to use for CCT [COMPASS]

Use for...	Population
Mechanism / CCT claims	mg10 min20 11_21 (+ optional high- T_j slice on top)
Sensitivity	mg10 min10
QC storytelling (“why we filter”)	FP, mg10 min0
Avoid for CCT claims	FP min0 alone — cameo noise

Part 6 — HAND17 Slide 5 — the geometry Ginger can picture

Slide 5 is titled something like “**NCAA team $\hat{A}_i$ interval overlap.**” Think of each team-season as a **horizontal bar** on a talent axis:

- **Left end** = weakest player’s talent on the roster
- **Right end** = strongest player’s talent

- **Bar height / position** = how good the team is overall (τ_j)

Facts from the slide (~2015–2019 window):

- $H_{\text{sort}} \approx 0.10$ — teams are **not** neatly sorted; almost everyone overlaps almost everyone.
- ~86% of the talent grid is covered by **more than one** team at once.
- At the **most common** talent level, ~3,117 **team-seasons** have *some* player there — the pond is **crowded horizontally**.
- **Typical roster span** $\approx 2.85 z$ — every team has weak and strong players; width does not collapse as teams get better.
- As τ_j rises, **intervals** shift right — contenders have higher floors **and** higher ceilings.

The Squid vs Jackal picture:

- Squid star at +2 z might be **near the top** of a mid-level team’s window — locally the big fish.
- Jackal player also at +2 z might sit on a team whose window runs to +2.5 or +3 — same number on the stat sheet, **more peers** at that level.

What Slide 5 shows: *where* the fish swim.

What it does not show: *who gets drafted* from each pond.

What BDP adds: draft-ecosystem teams sit in a **higher, tighter** τ_j band.

What CCT needs next: draft rate **on top of** the interval picture — at fixed $\hat{A}$.

Part 7 — What the model already taught us (so the empirics are not fishing)

You are not starting from zero. The **fake league** (Phase B) and the **MLE fit** already encode the Squid/Jackal logic.

7.1 The score and the knobs

Piece	Symbol	Role in CCT
Own ability	A_i ($\hat{A}_i$ in data)	How good is the fish?
Teammate pond (leave-one-out)	poolq_loo	How good are teammates excluding me? — best empirical congestion axis
Team congestion in score	L_C	Viable-peer crowding on roster (viability map)
Congestion weight	λ	How hard L_C penalizes ranking
Viability sharpness	γ, θ	Who counts as a “peer” for congestion
Selection rule	top-K, Gibbs	Who wins slots after scores
Sorting on assignment	ρ	Separate from score — do not merge

[SCOUT]: In Squid/Jackal prose, say “**teammates around me.**” In plots, use poolq_loo first. τ_j (team mean) **includes the player** in the average — slightly wrong label for “pond” (use as secondary axis only).

7.2 Phase B sim (HAND16) — the wind tunnel [COMPASS]

Old slider exercises still matter:

- $\lambda > 0$ in the **score** bends who gets picked **even when the winner rule is fixed** — pure mechanism proof that “standing out matters.”
- γ changes who counts as congestion-bearing.
- ρ (assignment sorting) moves team structure — separate from draft score.
- θ marks edge-case Squids and Jackals near the viability threshold.
- K/N — fraction of the pool that gets the scarce outcome (K winners / N). **Army captain promotion $\approx 35\text{--}40\%$** (high baseline); **NCAA draft $\approx 2\text{--}2.5\%$** on your panel (scarcer). Army **screams** because the **macro hero tail drop is robust**, not because Army K/N is smaller than NCAA. In the **fake league**, K/N is a **sim dial** ($\theta = F^{-1}(1-K/N)$; gallery often ~10%) — separate from either empirical rate.
- **PD20 Gibbs SELECT** — softens sharp top-K edges; sensitivity for near-miss cases.

NCAA $\rho \approx 0$ on the hero panel → **low sorting, high overlap** → exactly the regime where **within-team rank** might disambiguate same-z players better than league sorting does.

7.3 MLE fit — the panel “acts as if” congestion matters [SCOUT]

Bernoulli MLE on PPM (2013–2021 panel, γ fixed, refit λ and t):

- $\hat{\chi} \approx 2.6$
- $\hat{t} \approx 1.1$

Plain English: Even while the **marginal hero curve** looks flat at the elite tail, a fitted draft model that includes **congestion in the score** finds **non-trivial λ** — the data behave **as if** crowded rosters drag draft log-odds down **after controlling for ability**.

Critical nuance: MLE conditions on **A and L_C jointly**. The hero **marginal** curve (draft rate vs poolq_loo bin **only**) **averages over** ability levels. It can **hide** “same A, worse pond hurts.” That is why your CCT can be **true in the model layer** and **invisible on the old slide**.

Part 8 — The hero chart history (why July looked like you, and why we moved on)

Chart	What it showed	Does it prove CCT?
Original hero (R.I.P.) ~62k rows, min20, no mg10	Middle rise, dip in elite bin	Looked like CCT — but contaminated
POST-QC hero (mg10 min20)	Middle rise, flat elite tail	Does not prove or disprove — wrong averaging
OBPM/BPM Track C	Stronger rise, no elite dip	Still marginal; not inverted-U rescue
July mg=0 replay	Elite dip returns	Cameo/QC artifact — not the subtle signal

[SCOUT + COMPASS, joint verdict]:

- The July inverted-U tail was mostly who got dumped into bin 16 (fragmentary teams, cameo players) — not the congestion whisper you want to defend.
- The POST-QC flat tail means: if CCT is true in NCAA, it will not show up as “bin 16 dips” on a single marginal curve. You must hold $\hat{A}$ fixed and compare ponds.

That matches your instinct: Army screams; NCAA needs a microscope, not more ESPN seasons. [COMPASS] You and Alex already agree more coverage alone is unlikely to move the CCT needle.

Part 9 — What we have not proven yet (honest fence)

Marginal histograms and ventile hero curves answer:

- “What does the league look like?”
- “What do draft-history teams look like vs the league?”

They do not answer:

- “For player X at $z = +1.5$, does draft probability drop when we move them from a Squid-like team to a Jackal-like team?”

No disagreement between SCOUT and COMPASS on this.

Your conviction remains plausible and structurally supported. It is not yet extracted from NCAA data in the form you need for the dissertation defense — a conditional display where Squid and Jackal are explicitly compared at the same talent.

Part 10 — The microscope: what to plot next

[COMPASS] sequencing: BDP Act I is done. Act II should be one family of conditional figures before reopening MLE ordering or PD14 debates.

[SCOUT]: We are not running these until you pick them for the next HAND deck. This is the menu, in priority order.

Priority 1 — The direct Squid vs Jackal plot (both agents: start here)

Fix a narrow $\hat{A}_i$ band (e.g. z between 1.5 and 2.0, or top decile of ability). Within that band only, plot draft rate vs:

- poolq_loo ventiles (preferred congestion axis), or
- T_j ventiles (secondary)

Squid proxy: mid pond / mid T_j , same $\hat{A}$ band.

Jackal proxy: top pond / top T_j , same $\hat{A}$ band.

If CCT holds: the Jackal bar is lower than the Squid bar at same individual talent.

Filters: mg10 min20 11_21; PPM canonical; OBPM/BPM as robustness panels.

This is your thought experiment literally in the data.

Priority 2 — Two-way heatmap [SCOUT]

Draft rate in every cell ($\hat{A}_i$ ventile $\times$ poolq_loo ventile). The familiar hero curve is a column average — it smears the story. Look for lower rates in the upper-right (high ability, high pond) vs upper-middle (high ability, medium pond).

[SCOUT]: Signal may live only in the top 10–15% of $A \times$ top 10–15% of poolq_loo. Full-population plots drown it. Use wider bins at the top if cell counts get thin (~1,100 total drafts).

Priority 3 — Within-team rank [SCOUT] §11.3, [COMPASS] Step B

Marginal $\hat{A}$ and T_j miss the missing axis. For each player-season compute:

Construct	Meaning	CCT role
Roster percentile	Where $\hat{A}_i$ sits within the team (0 = bottom, 1 = top)	“Am I the big fish on this team?”
Headroom	Gap to best teammate, or count of peers within ϵ below	“How many rivals at my level?”
N_peers above θ	Teammates above a viability threshold near me	Empirical face of L_C

Plot: Draft rate vs roster percentile, T_j faceted by T_j quartile — four small panels. CCT signature: at high T_j , only the top of the team keeps draft mass; at mid T_j , same z but higher rank $\rightarrow$ higher rate.

Priority 4 — Interval overlap + draft [SCOUT]

Redo HAND17 Slide 5 logic on the hero panel: color drafted vs not drafted; top T_j quartile only; ask whether drafted players cluster on the **leading edge** of each team's interval.

Priority 5 — Model-aligned diagnostics [SCOUT]

- L_j^C vs T_j by ventile — high T_j should mean higher team congestion; if $\lambda > 0$, lower draft odds at given A.
- **MLE residuals** — observed minus predicted draft rate by poolq_loo at **fixed A** — shows where to zoom.

+DFT zoom (both agents)

Repeat Priority 1–2 on the **+DFT subsample only**. BDP sd compression hints the signal may **sharpen** inside the draft ecosystem. Use +DFT as **magnification**, not as “the answer.”

When you green-light: [SCOUT] build plan

Script name (locked): `pass_a_congestion_conditional.py`
Panel: hero `mg10 min20 11_21` — reuse existing `pass_a_empirical_bundle` path, **no new scrape**.
Outputs to `basic_data_plots/`:

- `BDP_draft_rate_by_ai_band_and_Tj.png` (or poolq_loo variant)
- `BDP_draft_rate_by_roster_percentile_and_Tj.png`
- `BDP_draft_rate_ai_x_poolq_loo_heatmap.png`
- JSON with **cell counts** (guard sparse bins)

Part 11 — What we will not do next (save your porch time)

[SCOUT] list, [COMPASS] endorses:

1. **More ESPN seasons alone** — unlikely to reveal CCT; you and Alex agree.
2. **Another marginal poolq_loo ventile run** hoping the elite bin dips — Track C closed that path on POST-QC data.
3. **Promoting OBPM or BPM to canonical hero** — useful appendix magnifier, not the main estimand.
4. **Confusing T_j with poolq_loo** in talk or slides — Squid/Jackal is about **teammates around the player**, not the team mean label.
5. **Another full ladder of marginal histograms** — Act I is complete.

Part 12 — One paragraph for Alex (Charles or Ginger can read this aloud on the porch)

[SCOUT] draft, [COMPASS] adopts verbatim:

On cleaned NCAA data the simple hero curve rises through the middle and doesn't show a robust elite dip — the July tail was mostly bad roster rows. But interval overlap shows elite teams stack several players at the same talent level where a mid-major might have one star. Our fitted draft model already puts weight on congestion (λ). The next empirical step is not more bins on the old chart — it's holding individual talent fixed and comparing draft rates on crowded versus uncrowded ponds. That's the Squid versus Jackals test.

Part 13 — Who does what when you pick the next plot

Job	Owner
BDP $\hat{A}$ / T_j / roster ladders	Done (you + scripts)
Matched $\hat{A}$ × pond draft-rate plots	SCOUT
Interval overlap conditional on draft	SCOUT
L_C vs T_j diagnostic	SCOUT
Slide wording, ladder order, sequencing	COMPASS
HAND deck assembly, Change Picture	Charles

When you name the first plot for the next HAND deck, tell us the slide title + filter ladder. We are ready.

Part 14 — Bottom line (the great news you waited for)

1. **Your CCT is coherent.** It fits Army, the model (λ , score vs select), HAND17 intervals, and today's BDP +DFT overlays.
2. **The BDP deck is a success.** You now have a shared visual language: FP → hero panel, three perf metrics, draft-ecosystem zoom. **Act I is closed.**
3. **Your OBPM/BPM compression instinct is real.** Comprehensive metrics **magnify** how alike draft-ecosystem teams look. Use as sensitivity, not canon.
4. **Congestion is not roster headcount** at min20. It is **talent stacked into similar-minute rosters** — rank and interval geometry.
5. **The old hero elite dip is not coming back** on POST-QC marginals — and that is **not** a defeat. CCT was never primarily “bin 16 goes down.”
6. **The way forward is conditional, not marginal:** same $\hat{A}$, compare draft odds across **pond thickness** and **within-team rank** — on **mg10 min20**, zoomed to the top of the ability × pond grid.
7. **SCOUT and COMPASS agree completely** on substance today. No fork in the road — one next step: **Priority 1 matched plot**.

Glossary (inline — no scavenger hunt)

Term	Plain English
$\hat{A}_i$	Empirical player ability (PPM/BPM/OBPM z within season)
$**T_j^{**}$	Mean player ability on team j that season
poolq_loo	Leave-one-out mean teammate quality — “pond excluding me”
L_C	Team congestion in the score (viability-weighted peers)
λ	Weight on congestion in ranking: $S = A - \lambda L_C$
DFT / +DFT	Team-seasons with ≥ 1 draftee in 2011–2021; orange player histogram = all players on those teams (not drafted-only)
mg10	Drop team-seasons with ≤ 10 games
min20	Players need ≥ 20 minutes that season
H_sort	How sorted teams are on talent (~0.1 = heavy overlap)
Hero	Draft rate vs teammate-quality ventiles (outcome chart)
CCT	Big fish / small pond vs small fish / big pond at same talent
K/N	Winners / pool size. Army promotion ~35–40%; NCAA draft ~2–2.5%. “Army screams” = macro tail drop, not tiny Army K/N

Where the pieces live (if you want to dig later)

Artifact	Path
BDP slide exports	HEROs_and_PASSes/slides/Basics_data_plots_HAND/
BDP source PNGs + JSON	HEROs_and_PASSes/basic_data_plots/
HAND17 interval slide	slides/CHAR_PD17_HAND/Slide5.png
Agent memo pair (technical)	SCOUT_and_COMPASS/20260820_*.md
Track C BPM/OBPM note	SCOUT_and_COMPASS/20260820_1302_SCOUT_to_COMPASS_track_c_bpm_obpm_robustness.md
Big-picture dissertation arc	HEROs_and_PASSes/PD20_22_campaign_big_picture.md
Binding score vs select	3-Master_Plan/BINDING_Selection_is_its_own_step.md

Charles — you asked your fellows **SCOUT** and **COMPASS** for a plain-spoken porch read you could share with **Ginger**. Here it is. The pond is real. The fish can be the same size. The question is how many same-size fish swim beside each one — and whether the draft notices. We believe you. When you’re ready, we’ll build the microscope.

(Twenty years married September 11 — congratulations to you and Ginger.)

— **SCOUT** and **COMPASS**, your fellas