

Tier 1 roadmap (SCOUT / VECTOR) — living checklist

Location: `sports/documents/tier_1_roadmap.md` (alongside other sports-facing notes; not a duplicate of theory memos in `5-Manuscript`.)

Notebook: `sports/535_sports_tier_1.ipynb`

Purpose: Single place for the mechanical pipeline, the `df` column contract after each stage, and **CELL 0** switches. Update this file as CELL 4+ and modeling land.

Contract (what “must be true” for implementation): this file + the **module docstrings / function contracts** in `sports_pipeline/tier1_mechanism_vars.py` and `sports_pipeline/panel_build.py` (especially LOO / `poolq_loo` behavior). Manuscripts in `5-Manuscript` are **design intent and exposition**, not the execution contract unless we explicitly mirror a rule here.

Related theory (`5-Manuscript` , stay at source — link, don't copy):

- `Vector_to_Scout_Tier1_Modeling_Direction.md` — main Tier 1 modeling direction.
- `tier_1_model.md` — model layer notes.
- `2026_0430_Paper7_feedback.md` — theory vs minimal model; scarcity.

Optional depth (broader VECTOR notes): `Vector_Master_Theory_and_Modeling_Notes*.md` , `Vector_Questions_and_Modeling_Thoughts.md` , `Vector_Evans_Reaction_and_Theoretical_Expansion_Notes.md` — pull excerpts into this roadmap only when a line becomes a **coded rule** or sample definition.

Pipeline (numbered)

1. **Load the panel** — `load_panel(CFG)` → `df` from `player_season_panel_530.csv` at **workspace root** `datasets/mbb/` (via `sports_pipeline.paths.panel_530_csv()` , same as **530 conductor**).
2. **Set `perf` from a metric** — e.g. PPM via `assign_perf_from_metric(df, PERF_METRIC)` with `PERF_METRIC` defined in **CELL 0**.
3. **Recompute legacy `poolq_loo` (and `poolq_sq`)** — `recompute_teammate_loo_pool_quality(...)` .

Why recompute instead of trusting the CSV?

- LOO pool quality must be built off the `perf` **you just chose**; a saved `poolq_loo` may reflect an older or different `perf` definition.
- Optional **winsor** on `poolq_loo` via `CFG.poolq_winsor_quantiles` is applied in this step when set.

Comparison: `poolq_loo` (legacy) and `congestion_quality` (Tier 1) can differ slightly when some roster rows have missing `perf` ; the Tier 1 helper uses only rows with valid `perf` in the `(team_id, season)` group for denominators. See `sports_pipeline/tier1_mechanism_vars.py` module docstring.

4. **Add Tier 1 mechanism columns** — `add_tier1_mechanism_variables(...)` :
 - `congestion_quality` — LOO **mean** teammate `perf` (valid Q for Tier 1).
 - `congestion_crowding` — LOO **sum** teammate `perf` .
 - `congestion_crowding_weighted` — optional weighted variant (minutes × `perf` logic); gated by **CELL 0** `COMPUTE_WEIGHTED_CROWDING` .

- If `MIN_MINUTES_TIER1` is set in **CELL 0**, rows below that **minutes** threshold are dropped **before** group stats.

5. **CELL 4 — Tier 1 EDA** (notebook): On a sample with valid `PRIMARY_POOL_MODE` pool column, `congestion_crowding`, `minutes`, `Y_draft`: descriptive `describe()`; **binned** mean draft vs Q using `assign_poolq_bin_labels` with `CFG.ventiles` and `CFG.poolq_binning` (530-style); numpy LPM $Y \sim 1 + Q + Q^2$; multivariate linear $Y \sim 1 + Q + Q^2 + C + \text{minutes}$ when Q is quality (if Q is crowding, **omit duplicate C**); **PD-style** predicted curve over a Q grid with C and minutes at sample medians (linear approx, clipped to [0,1]).

`df` column contract (Tier 1–relevant)

After steps 1–3 (after notebook CELL 2):

- Required for assertions in the notebook: `Y_draft`, `poolq_loo`, `minutes`.
- Always present after recompute (when inputs allow): `perf`, `poolq_loo`, `poolq_sq`.
- Plus **all other columns** from the 530 CSV (identifiers such as `athlete_id`, `team_id`, `season`, and any merges from 530).

After step 4 (after notebook CELL 3):

- Same as above, **plus**: `congestion_quality`, `congestion_crowding`, and usually `congestion_crowding_weighted` (or NaNs if weighted path is off / invalid inputs).
-

CELL 0 — variables / boolean switches (current)

Kind	Name	Role
Run gate	<code>RUN_CELL1</code>	Environment · imports · panel path
Run gate	<code>RUN_CELL2</code>	Load panel · PPM · legacy <code>poolq_loo</code>
Run gate	<code>RUN_CELL3</code>	Tier 1 mechanism variables
Run gate	<code>RUN_CELL4</code>	Tier 1 EDA (bins, LPM, PDP-style Q grid)
Knob	<code>PERF_METRIC</code>	Passed to <code>assign_perf_from_metric</code> (default <code>"ppm"</code>)
Knob	<code>MIN_MINUTES_TIER1</code>	Optional row filter before Tier 1 group stats; <code>None</code> = full sample
Knob	<code>PRIMARY_POOL_MODE</code>	<code>"quality"</code> vs <code>"crowding"</code> for <code>tier1_primary_pool_column</code> (regressions later)
Knob	<code>COMPUTE_WEIGHTED_CROWDING</code>	Whether to build <code>congestion_crowding_weighted</code>

Pattern matches `540_tenure_pipeline.ipynb` CELL 0: one switchboard, re-run after changes.

Next (placeholder)

- CELL 5+: proper logit / `statsmodels` ; cluster SEs; season or team FE; richer PDP; export figures to `CFG.exports_dir` if you want parity with 530 PNG exports.

Changelog

Date	Note
2026-05-06	CELL 4 in 535: binned draft vs Q, quadratic LPM, C + minutes controls, PDP-style Q grid; <code>RUN_CELL4</code> in CELL 0.
2026-05-06	<code>paths.project_root</code> = Ivy_Net workspace root so <code>datasets/mbb/</code> matches 530 output (not <code>sports/datasets/mbb/</code>).
2026-05-05	File lives under <code>sports/documents/</code> ; contract vs theory section + links to <code>5-Manuscript</code> .
2026-03-31	Initial file: pipeline steps, recompute rationale, <code>df</code> contract, CELL 0 table.