VA Weekend Checklist — Tenure Pipeline

Goal: Get pipeline fully up to speed before introducing new URLs from `faculty_url_suggestions.csv`

⚡ TONIGHT — Do These Two Things Before You Log Off

> *Cell 0 flags are **already set correctly** — do NOT change them. Just commit and launch.*

1. Commit everything in Cursor Source Control

- Stage all → paste this message → commit → push:

```
cleanup: remove current_documents symlink deps, fix broken notebook path + Cell 0 flags for overnight run
```

2. Launch both Slurm jobs from the terminal

```
cd ~/Ivy_Net
sbatch build_openalex_cache.slurm   # overnight cache builder — runs alone, no conflicts
sbatch pipe_job.slurm               # fills HTML gaps + rebuilds panel + runs analysis
```

3. Verify both jobs are queued

```
squeue -u dzk3ja
# Should see TWO jobs listed
```

4. Log off — both jobs run unattended overnight

What runs tonight in pipe_job: - Cell 3 Download → fills 553 missing HTML files - Cell 3 Retry → retries any CDX failures (CDX discovery is OFF — quota protected) - Cell 4 → re-parses all HTML - Cell 5 → rebuilds longitudinal panel - Cell 7 → enriched annual panel (GAP_TOLERANCE = 4) - Cell 8 → LOO peer pool metrics - Cell 9 → inverted-U plot

What does NOT run (intentional): - Cell 6A — author ID resolution (not needed yet) - Cell 6B — OpenAlex works fetch — OFF to avoid conflict with cache builder writing same file

---

☀️ TOMORROW MORNING — Check Results

Check jobs finished

```
squeue -u dzk3ja         # should be empty if both finished
ls -lh ~/Ivy_Net/slurm_out/*.out   # check log files exist
tail -50 ~/Ivy_Net/slurm_out/slurm-pipe_job-*.out  # read pipe_job summary
tail -20 ~/Ivy_Net/slurm_out/slurm-oa_cache-*.out  # read cache builder summary
```

If pipe_job finished — check the output

Look for in the log: - `Panel rows` count (should be > previous run) - `Unique faculty` count - Stage 9 inverted-U plot saved - No ERROR lines

If cache builder finished — run Cell 6B

```
# In the notebook Cell 0, change:
#   RUN_CELL6B = False  →  RUN_CELL6B = True
#   (all other flags: set to False except 6B)
cd ~/Ivy_Net && sbatch pipe_job.slurm
```

📋 NEXT STEPS (After overnight jobs) — In Order

Phase 3 — Expand Pilot to Top-40 Schools

Step A — Resolve author IDs for schools 21–40 (Cell 6A) - [ ] In Cell 0: set `STAGE6_PILOT_TOP_N = 40` - [ ] Set `RUN_CELL6A = True`, all others `False` - [ ] `sbatch pipe_job.slurm` - [ ] Confirm: author ID count grows from 15,519

Step B — Fetch works by year for new authors (Cell 6B) - [ ] In Cell 0: set `RUN_CELL6B = True`, all others `False` - [ ] `sbatch pipe_job.slurm` - [ ] Confirm: works file grows; snapshot cache used where available

Step C — Re-run full analysis - [ ] In Cell 0: set `RUN_CELL7 = True`, `RUN_CELL8 = True`, `RUN_CELL9 = True`, all others `False` - [ ] `sbatch pipe_job.slurm` - [ ] Compare Stage 9 inverted-U plot to previous run

---

Phase 4 — Review URL Suggestions (Mac Local)

Do this AFTER Phase 3 is complete and results look stable.

Step D — rsync results back to Mac

```
# From your Mac terminal:
./scripts/rsync_pull_from_hpc.sh
```

Step E — Review suggestions CSV - [ ] Open: `tenure/tenure_pipeline/faculty_url_suggestions.csv` - [ ] Sort by `mean_recs` ascending (worst schools first) - [ ] For each school: check `suggested_url` — does it look like a **CS faculty page**? - [ ] **Red flags:** generic university staff directories, aging/CTE/HR pages - [ ] **Green flags:** subdomain has `cs`, `eecs`, `cse`, `computing`; path has `faculty`, `people` - [ ] Promising leads: - Rice University (score 97, 2000–2024) — `rice.edu/Computer/facultystaff.html` - Kent State (score 84, 2019–2025) — `kent.edu/cs/faculty-staff-directory` - Lehigh (score 83, 2014–2026) — `lehigh.edu/academics/faculty` - Northwestern (score 83, 2012–2024) — check if it's the CS directory

Step F — Add good URLs to the pipeline

```
# Edit tenure/tenure_pipeline/r1_cs_departments.csv
# Then re-run: Cell 2 → Cell 3A → Cell 3B → Cell 4 → Cell 5
```

Git Sync Rule — After Every sbatch + Results Look Good

```
cd ~/Ivy_Net
git add -A
git status        # verify: no *.jsonl, no *.html, no slurm-*.out / slurm_out noise staged
git commit -m "Weekend run: <brief description>"
git push
```

Quick Reference — Key Commands

```
# Check queue
squeue -u dzk3ja

# Monitor a job live
tail -f ~/Ivy_Net/slurm_out/slurm-<JOBNAME>-<JOBID>.out

# Submit jobs
cd ~/Ivy_Net && sbatch pipe_job.slurm
cd ~/Ivy_Net && sbatch build_openalex_cache.slurm

# Pull HPC outputs to Mac (what Git ignores — see ../../scripts/DATA_SYNC.md)
./scripts/rsync_pull_recent_hpc.sh quick   # slurm_out + sweep; use `all` to add tenure/tenure_pipeline
./scripts/rsync_pull_from_hpc.sh           # default: tenure/tenure_pipeline only

# Push code to Rivanna (run from Mac terminal)
./scripts/rsync_push_to_hpc.sh
```

Workspace Notes

Workspace	When to Use
`Rivanna.code-workspace`	Slurm jobs, HPC data, remote debugging
`Ivy_Net.code-workspace`	Code edits, git, rsync, local runs

Explorer shows 3 roots: `Ivy_Net`, `tenure`, `sports` — notebooks are directly visible, no symlinks.

Last updated: 2026-04-15 — overnight run plan + for-dummies rewrite